

```

from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.tree import DecisionTreeClassifier
import pandas as pd

```

```

df=pd.read_csv("./dataset/train_clean_data.csv",index_col=0)
df.head()

```

	Loan_ID	Dependents	ApplicantIncome	CoapplicantIncome	LoanAmount
1	LP001003	1	4583	1508.0	128.0
2	LP001005	0	3000	0.0	66.0
3	LP001006	0	2583	2358.0	120.0
4	LP001008	0	6000	0.0	141.0
5	LP001011	2	5417	4196.0	267.0

	Loan_Amount_Term	Credit_History	gender_Male	married_Yes
1	360.0	1.0	1	1
2	360.0	1.0	1	1
3	360.0	1.0	1	1
4	360.0	1.0	1	0
5	360.0	1.0	1	1

	education_Not Graduate	property_area_Semiurban
1	0	0
2	0	0
3	1	0
4	0	0
5	0	0

	self_employed_Yes	Loan_status_Y
1	0	0
2	1	1
3	0	1
4	0	1
5	1	1

```

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 517 entries, 1 to 613
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Loan_ID                               517 non-null    object
1   Dependents                           517 non-null    object
2   ApplicantIncome                       517 non-null    int64
3   CoapplicantIncome                     517 non-null    float64
4   LoanAmount                            517 non-null    float64
5   Loan_Amount_Term                       517 non-null    float64
6   Credit_History                         517 non-null    float64
7   gender_Male                           517 non-null    int64
8   married_Yes                           517 non-null    int64
9   education_Not Graduate                 517 non-null    int64
10  property_area_Semiurban                517 non-null    int64
11  property_area_Urban                    517 non-null    int64
12  self_employed_Yes                      517 non-null    int64
13  Loan_status_Y                          517 non-null    int64
dtypes: float64(4), int64(8), object(2)
memory usage: 60.6+ KB

df.shape

(517, 14)

df['Loan_status_Y'].value_counts()

Loan_status_Y
1    360
0    157
Name: count, dtype: int64

''' Train Test Split'''
X=df.drop(['Loan_ID','Dependents','Loan_status_Y'],axis=1)
y=df['Loan_status_Y']

from sklearn.model_selection import train_test_split
X_train,X_test, y_train,y_test=train_test_split(X,y,
test_size=0.33,stratify=y,random_state=42)

!pip install --upgrade scikit-learn
!pip install --upgrade imbalanced-learn

Requirement already satisfied: scikit-learn in c:\users\hp\anaconda3\
lib\site-packages (1.6.1)
Requirement already satisfied: numpy>=1.19.5 in c:\users\hp\anaconda3\
lib\site-packages (from scikit-learn) (1.26.4)
Requirement already satisfied: scipy>=1.6.0 in c:\users\hp\anaconda3\
lib\site-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in c:\users\hp\anaconda3\

```

```
lib\site-packages (from scikit-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\hp\
anaconda3\lib\site-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: imbalanced-learn in c:\users\hp\
anaconda3\lib\site-packages (0.13.0)
Requirement already satisfied: numpy<3,>=1.24.3 in c:\users\hp\
anaconda3\lib\site-packages (from imbalanced-learn) (1.26.4)
Requirement already satisfied: scipy<2,>=1.10.1 in c:\users\hp\
anaconda3\lib\site-packages (from imbalanced-learn) (1.13.1)
Requirement already satisfied: scikit-learn<2,>=1.3.2 in c:\users\hp\
anaconda3\lib\site-packages (from imbalanced-learn) (1.6.1)
Requirement already satisfied: sklearn-compat<1,>=0.1 in c:\users\hp\
anaconda3\lib\site-packages (from imbalanced-learn) (0.1.3)
Requirement already satisfied: joblib<2,>=1.1.1 in c:\users\hp\
anaconda3\lib\site-packages (from imbalanced-learn) (1.2.0)
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in c:\users\hp\
anaconda3\lib\site-packages (from imbalanced-learn) (3.5.0)
```

```
from imblearn.over_sampling import SMOTE
smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

model1=DecisionTreeClassifier(class_weight='balanced')
model2=GaussianNB()
model3=LogisticRegression(solver='liblinear',max_iter=200,class_weight
='balanced')

model1.fit(X_train_res,y_train_res)
model2.fit(X_train_res,y_train_res)
model3.fit(X_train_res,y_train_res)

LogisticRegression(class_weight='balanced', max_iter=200,
solver='liblinear')

pred1=model1.predict_proba(X_test)
pred2=model2.predict_proba(X_test)
pred3=model3.predict_proba(X_test)

final_pred=(pred1+pred2+pred3)/3

final_pred
array([[0.39223731, 0.60776269],
       [0.40659715, 0.59340285],
       [0.02848186, 0.97151814],
       [0.04511311, 0.95488689],
       [0.35097035, 0.64902965],
       [0.23817243, 0.76182757],
       [0.29479918, 0.70520082],
       [0.10438856, 0.89561144],
       [0.47694483, 0.52305517],
```

[0.09476933, 0.90523067],
[0.13907493, 0.86092507],
[0.14761552, 0.85238448],
[0.06334219, 0.93665781],
[0.11690038, 0.88309962],
[0.10704854, 0.89295146],
[0.52204028, 0.47795972],
[0.05678688, 0.94321312],
[0.04764203, 0.95235797],
[0.97920344, 0.02079656],
[0.07782672, 0.92217328],
[0.77405891, 0.22594109],
[0.59550243, 0.40449757],
[0.03753762, 0.96246238],
[0.0702189 , 0.9297811],
[0.62788621, 0.37211379],
[0.40903588, 0.59096412],
[0.02358319, 0.97641681],
[0.16701355, 0.83298645],
[0.09439583, 0.90560417],
[0.61470569, 0.38529431],
[0.19094428, 0.80905572],
[0.19994528, 0.80005472],
[0.03651946, 0.96348054],
[0.94172623, 0.05827377],
[0.34386 , 0.65614],
[0.17219798, 0.82780202],
[0.64870456, 0.35129544],
[0.24293707, 0.75706293],
[0.18219011, 0.81780989],
[0.98506099, 0.01493901],
[0.60467578, 0.39532422],
[0.06690004, 0.93309996],
[0.85849533, 0.14150467],
[0.05396406, 0.94603594],
[0.30974308, 0.69025692],
[0.36173541, 0.63826459],
[0.37874488, 0.62125512],
[0.56636961, 0.43363039],
[0.3878724 , 0.6121276],
[0.10266172, 0.89733828],
[0.8235694 , 0.1764306],
[0.24950517, 0.75049483],
[0.48954473, 0.51045527],
[0.23865234, 0.76134766],
[0.57781499, 0.42218501],
[0.36480475, 0.63519525],
[0.10416488, 0.89583512],
[0.24594933, 0.75405067],

[0.06270525, 0.93729475],
[0.37853547, 0.62146453],
[0.08039977, 0.91960023],
[0.12339853, 0.87660147],
[0.24639923, 0.75360077],
[0.16321457, 0.83678543],
[0.14226174, 0.85773826],
[0.55648344, 0.44351656],
[0.2484255 , 0.7515745],
[0.10489714, 0.89510286],
[0.98770409, 0.01229591],
[0.97427746, 0.02572254],
[0.93081756, 0.06918244],
[0.98136976, 0.01863024],
[0.13181808, 0.86818192],
[0.05257288, 0.94742712],
[0.14986857, 0.85013143],
[0.11732665, 0.88267335],
[0.93969747, 0.06030253],
[0.15023051, 0.84976949],
[0.11413457, 0.88586543],
[0.42017344, 0.57982656],
[0.14599222, 0.85400778],
[0.96365373, 0.03634627],
[0.37458732, 0.62541268],
[0.77455218, 0.22544782],
[0.23732365, 0.76267635],
[0.96825252, 0.03174748],
[0.13666245, 0.86333755],
[0.12635167, 0.87364833],
[0.99196877, 0.00803123],
[0.97274661, 0.02725339],
[0.04683254, 0.95316746],
[0.94161224, 0.05838776],
[0.06478955, 0.93521045],
[0.05562148, 0.94437852],
[0.14495654, 0.85504346],
[0.26132345, 0.73867655],
[0.18902338, 0.81097662],
[0.12724161, 0.87275839],
[0.34600846, 0.65399154],
[0.09062174, 0.90937826],
[0.13865821, 0.86134179],
[0.39696446, 0.60303554],
[0.3783675 , 0.6216325],
[0.05767018, 0.94232982],
[0.13440323, 0.86559677],
[0.54936718, 0.45063282],
[0.16338474, 0.83661526],

[0.24888678, 0.75111322],
[0.04649542, 0.95350458],
[0.05810475, 0.94189525],
[0.06726684, 0.93273316],
[0.0347083 , 0.9652917],
[0.07277309, 0.92722691],
[0.04736222, 0.95263778],
[0.26392414, 0.73607586],
[0.13004317, 0.86995683],
[0.0616286 , 0.9383714],
[0.057139 , 0.942861],
[0.24471376, 0.75528624],
[0.28269427, 0.71730573],
[0.12476753, 0.87523247],
[0.24762113, 0.75237887],
[0.17149273, 0.82850727],
[0.08219659, 0.91780341],
[0.27276099, 0.72723901],
[0.07110289, 0.92889711],
[0.46447266, 0.53552734],
[0.13523122, 0.86476878],
[0.06061216, 0.93938784],
[0.18129636, 0.81870364],
[0.95369512, 0.04630488],
[0.98792005, 0.01207995],
[0.061814 , 0.938186],
[0.14926299, 0.85073701],
[0.48353623, 0.51646377],
[0.30743803, 0.69256197],
[0.65056053, 0.34943947],
[0.98192255, 0.01807745],
[0.24479796, 0.75520204],
[0.98449143, 0.01550857],
[0.14107747, 0.85892253],
[0.06430535, 0.93569465],
[0.98477633, 0.01522367],
[0.92251243, 0.07748757],
[0.99508196, 0.00491804],
[0.11150255, 0.88849745],
[0.03620421, 0.96379579],
[0.14602632, 0.85397368],
[0.07531721, 0.92468279],
[0.43218168, 0.56781832],
[0.20282846, 0.79717154],
[0.54620235, 0.45379765],
[0.3055709 , 0.6944291],
[0.96577597, 0.03422403],
[0.30581131, 0.69418869],
[0.98497405, 0.01502595],

```
[0.97204098, 0.02795902],
[0.21292266, 0.78707734],
[0.28882854, 0.71117146],
[0.96167384, 0.03832616],
[0.2462701 , 0.7537299 ],
[0.15338226, 0.84661774],
[0.27216106, 0.72783894],
[0.40335195, 0.59664805],
[0.23020923, 0.76979077],
[0.96751815, 0.03248185],
[0.97383931, 0.02616069],
[0.18876619, 0.81123381],
[0.47191153, 0.52808847],
[0.24653702, 0.75346298],
[0.9771367 , 0.0228633 ]])
```

```
model1.classes_
```

```
array([0, 1], dtype=int64)
```

Voting Classifier

```
clf1=DecisionTreeClassifier(class_weight='balanced')
clf2=GaussianNB()
clf3=LogisticRegression(solver='liblinear',max_iter=200,class_weight='balanced')

ecclf1=VotingClassifier(estimators=[('DT',clf1),('GNB',clf2),
('LR',clf3)], voting='hard')
ecclf1=ecclf1.fit(X_train_res,y_train_res)

predictions=ecclf1.predict(X_test)

predictions
array([0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0,
1,
1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
1,
1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,
0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1,
1,
0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1,
1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0,
0,
1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1,
0,
0,
```

```
1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0],
dtype=int64)
```

```
from sklearn.metrics import classification_report , confusion_matrix
print(classification_report(y_test,predictions))
print(confusion_matrix(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.70	0.62	0.65	52
1	0.84	0.88	0.86	119
accuracy			0.80	171
macro avg	0.77	0.75	0.76	171
weighted avg	0.80	0.80	0.80	171

```
[[ 32  20]
 [ 14 105]]
```

```
eclf2=VotingClassifier(estimators=[('DT',clf1),('GNB',clf2),
('LR',clf3)], voting='soft')
eclf2=eclf2.fit(X_train,y_train)
```

```
prediction=eclf2.predict(X_test)
```

```
from sklearn.metrics import classification_report , confusion_matrix
print(classification_report(y_test,prediction))
print(confusion_matrix(y_test,prediction))
```

	precision	recall	f1-score	support
0	0.67	0.60	0.63	52
1	0.83	0.87	0.85	119
accuracy			0.79	171
macro avg	0.75	0.74	0.74	171
weighted avg	0.78	0.79	0.79	171

```
[[ 31  21]
 [ 15 104]]
```

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
base_model1=DecisionTreeClassifier()
bagging = BaggingClassifier(estimator=base_model1, n_estimators=200,
random_state=0)
bagging.fit(X_train_res, y_train_res)
y_pred_bagging = bagging.predict(X_test)
print("Bagging Accuracy:", accuracy_score(y_test, y_pred_bagging))
```


Bagging Accuracy: 0.7660818713450293

```
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=100, random_state=0)
rf.fit(X_train_res, y_train_res)
y_pred_rf = rf.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
```

Random Forest Accuracy: 0.7953216374269005

```
from sklearn.ensemble import AdaBoostClassifier

boosting = AdaBoostClassifier(n_estimators=200, random_state=0)
boosting.fit(X_train_res, y_train_res)
y_pred_boost = boosting.predict(X_test)
print("Boosting (AdaBoost) Accuracy:", accuracy_score(y_test,
y_pred_boost))
```

Boosting (AdaBoost) Accuracy: 0.7777777777777778